

# Why use just one computer? (An intro to Distributed Systems)

João Mota



# On the menu

- I've got a system. Is it distributed?
- How processes communicate
- Multicast
- Hands-on!
- Final thoughts

# I've got a system. Is it distributed?

Symptoms of a distributed system:

- Your system has **more than one process**
- Your system's processes (can) **run on different computers**
- Your system's processes **communicate with each other** (via messages)

# Distributed systems in the wild

- The web and the internet
- E-mail
- ATM networks
- P2P applications (bittorrent and others)
- Swarm robotics
- IoT
- Autonomous driving
- Blockchain
- ...



# How processes communicate

In distributed systems, **communication is key**. Different computers can't really cooperate if they don't understand how to talk to each other.

Application	← Application specific protocol
Transport	← Communication between two processes
Network	← Communication between computers that <b>aren't</b> directly connected
Data Link	← Communication between two computers that <b>are</b> directly connected

# How processes communicate - Transport Layer

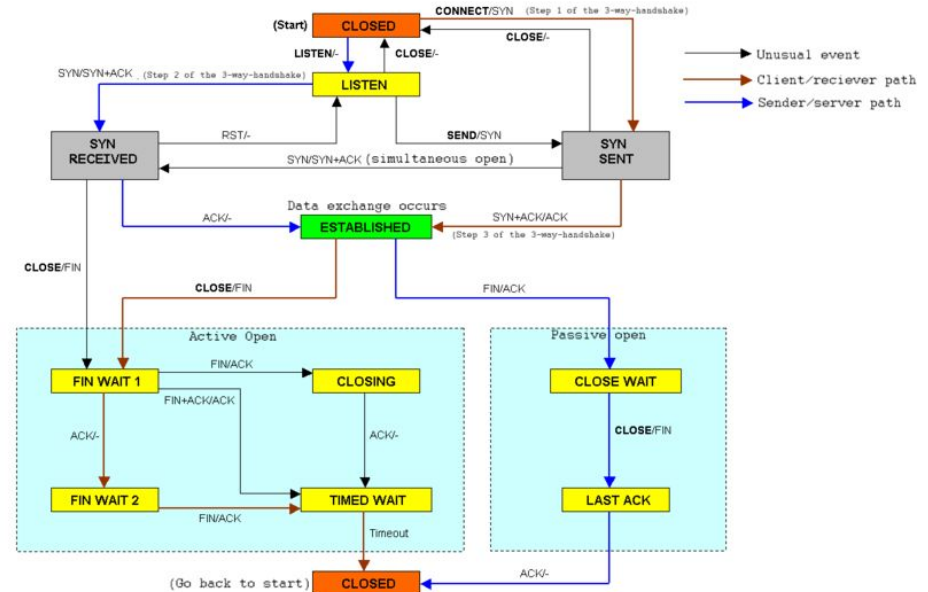
The two big transport layer protocols:

	UDP	TCP
Connection setup	No	Yes
Order	No	Yes
Reliability	No	Yes
Flow Control	No	Yes
Endpoints	*	1

# How processes communicate - Transport Layer

## Connection setup:

"Hi, I'd like to hear a TCP joke."  
"Hello, would you like to hear a TCP joke?"  
"Yes, I'd like to hear a TCP joke."  
"OK, I'll tell you a TCP joke."  
"Ok, I will hear a TCP joke."  
"Are you ready to hear a TCP joke?"  
"Yes, I am ready to hear a TCP joke."  
"Ok, I am about to send the TCP joke. It will last 10 seconds, it has two characters, it does not have a setting, it ends with a punchline."  
"Ok, I am ready to get your TCP joke that will last 10 seconds, has two characters, does not have an explicit setting, and ends with a punchline."  
"I'm sorry, your connection has timed out. ...  
Hello, would you like to hear a TCP joke?"



# How processes communicate - Transport Layer

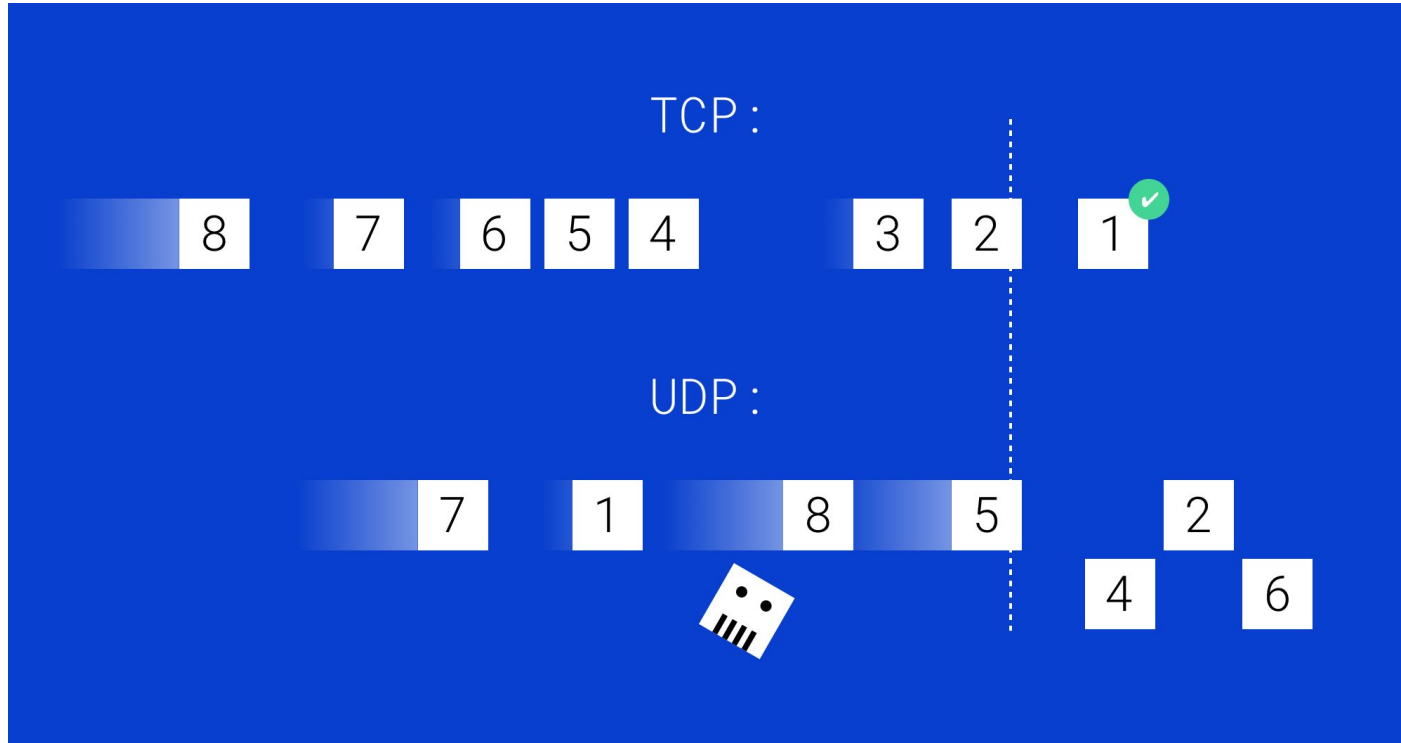
## Connection Setup - Advantages:

- More information = more control
  - Can ensure flow control
  - Can ensure order
  - Can ensure delivery
- Provides some security



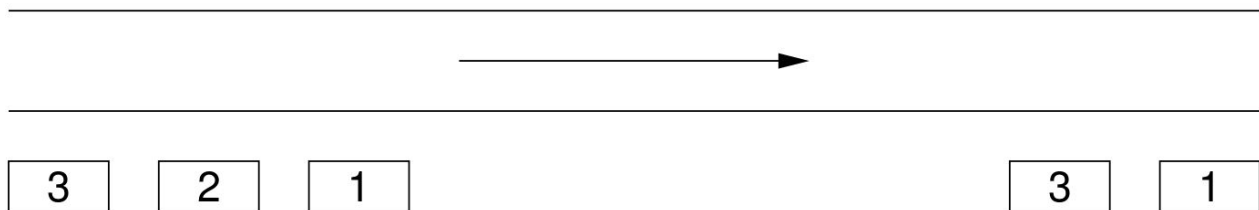
# How processes communicate - Transport Layer

Order:

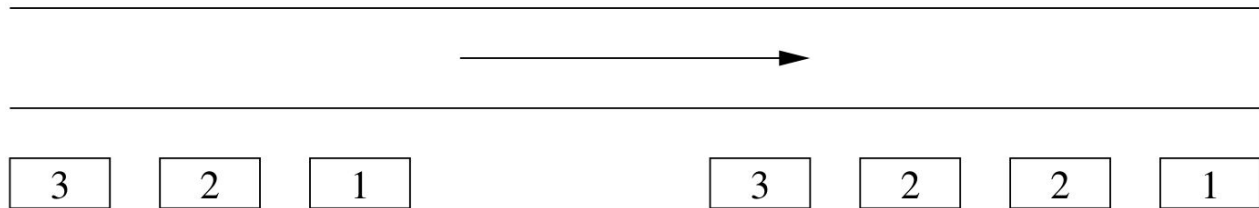


# How processes communicate - Transport Layer

Reliability - Packet loss:

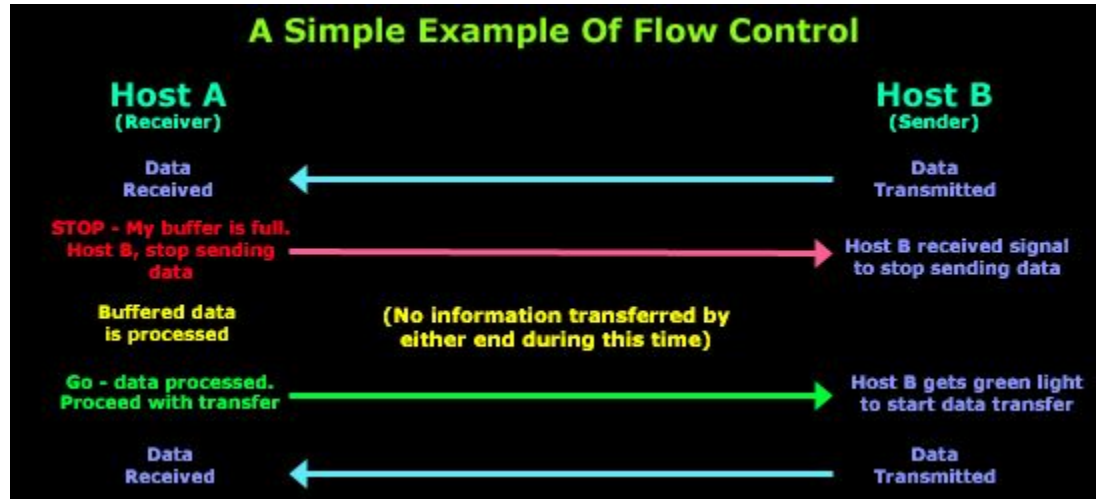


Reliability - Packet duplication:



# How processes communicate - Transport Layer

Flow control:



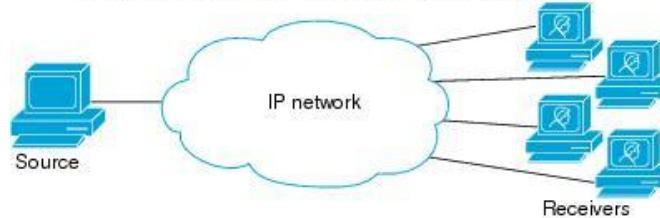
# How processes communicate - Transport Layer

Endpoints:

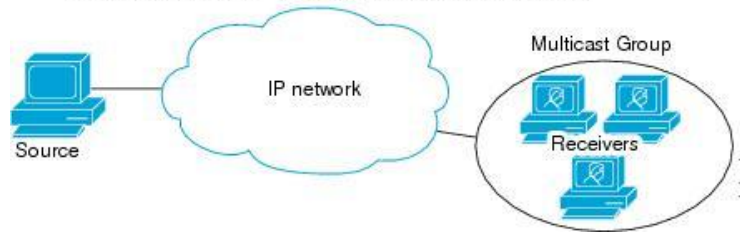
Unicast transmission—One host sends and the other receives.



Broadcast transmission—One sender to all receivers.




Multicast transmission—One sender to a group of receivers.



# How processes communicate - Transport Layer

How to choose the transport protocol for your application:



# How processes communicate - Application Layer

Here is where you can really define your system's way of communicating. You can design your protocol any way you want, but there are some useful frameworks:

- Binary
- Text
- JSON
- XML

# Multicast

- IP Multicast
- Application-level Multicast

# Multicast

Multicast: A naïve implementation

- N senders
- M receivers
- $N \times M$  single-ended channels
- M messages must be sent for each message
- Same message will probably traverse the same link multiple times
- The sender needs to know all of its receivers



# Multicast

## IP Multicast

- Use a spanning tree containing the sender, all receivers and all nodes in between them
- The sender does everything as if it were doing unicast
- Receivers **subscribe** to the multicast group to be added to the spanning tree
- They can also **unsubscribe** to be removed
- IP multicast can only be used within the same network and using UDP
  - Reliable multicast is **hard**
  - ISPs don't provide multicast cross-networks
  - IPv6 supposedly does!

# Multicast

## Application-level Multicast

- Build an overlay network
- Can use spanning tree
- Alternatively use epidemic algorithms

Hands-on!

[goo.gl/QLrdEk](https://goo.gl/QLrdEk)

# Final thoughts

So, why use just one computer?

- It's simpler
- You don't need your processes to be on different computers
- You don't even have multiple processes
- Distributed systems are a pain in the a\*\*!

# Multicast

Why are distributed systems even a thing?


- They enable communication (potentially) without spatial boundaries
- They enable collaborative work across multiple systems
- They allow for bigger systems
- Spreading functionality may be useful!
  - Redundancy
  - Sharding
  - Divide-and-conquer algorithms
  - Partial anonymity

# What now?

This workshop was just the tip of the iceberg.

- Remote Method Invocation (RMI)
- Cryptography
- Name resolution
- Serverless features
- Synchronization
- Fault Tolerance
- Consensus
- Atomic commitment

## DISTRIBUTED SYSTEMS



Maarten van Steen  
Andrew S. Tanenbaum

THIRD EDITION - VERSION 01